

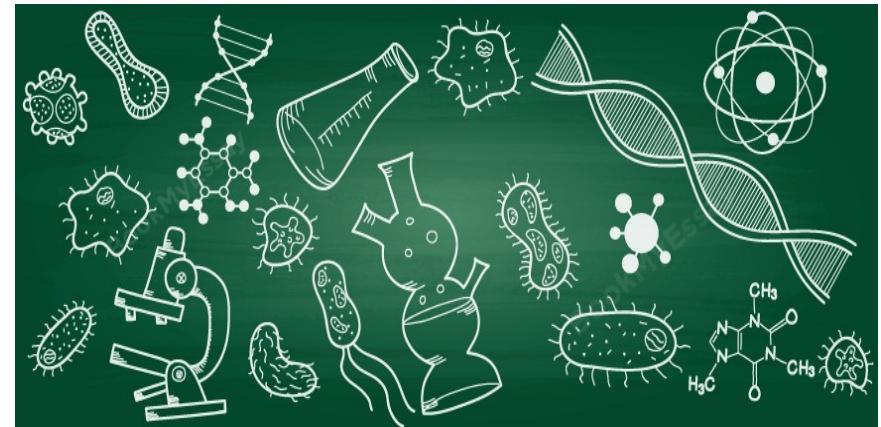
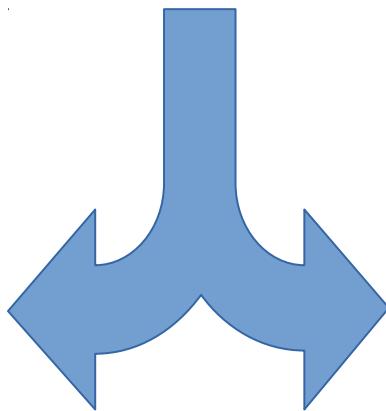
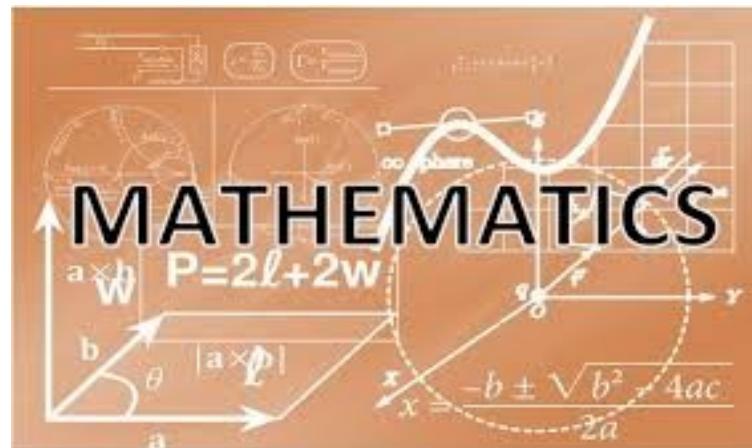
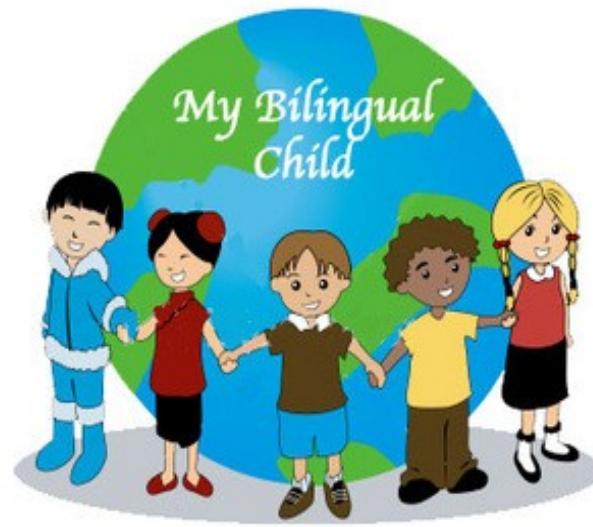
Bayesian Deep Learning Applications in Biomedicine

BayesLund 2019

Nikolay Oskolkov, NBIS SciLifeLab
Lund, 07.05.2019



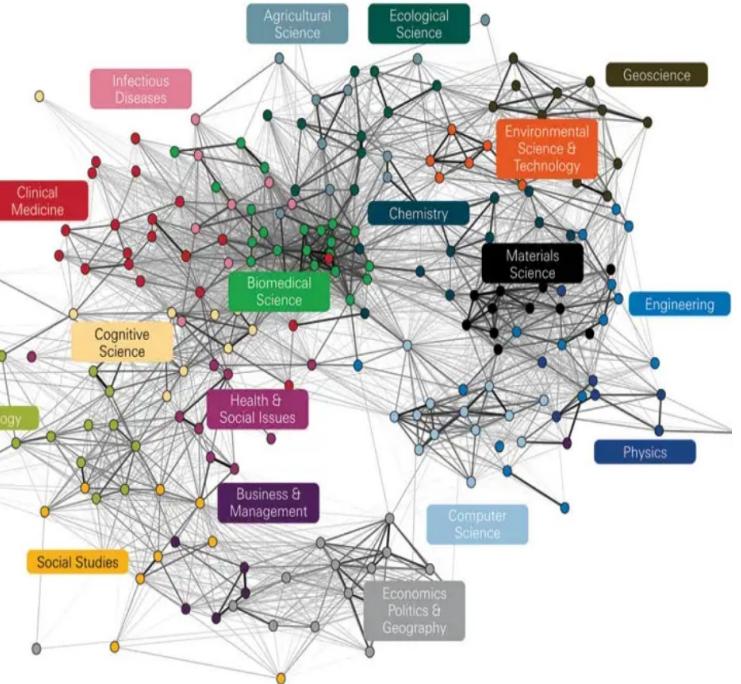
Bioinformatics in the Bayesian way



Diverse Audience



Different Sciences



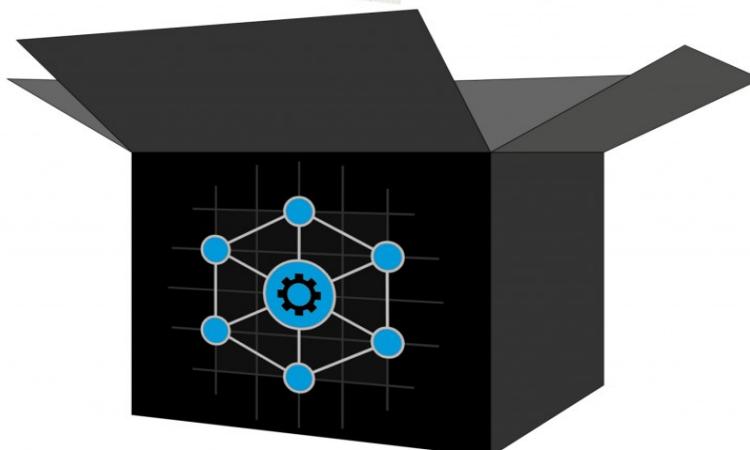
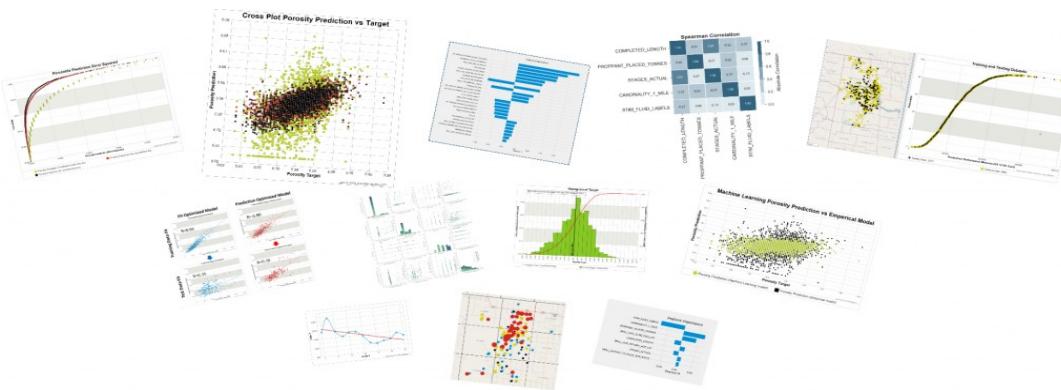
Machine Learning for Everyone





Machine Learning has nothing
to do with machines

... but mathematics



Machine Learning is not
a black box

... because it selects biomarkers

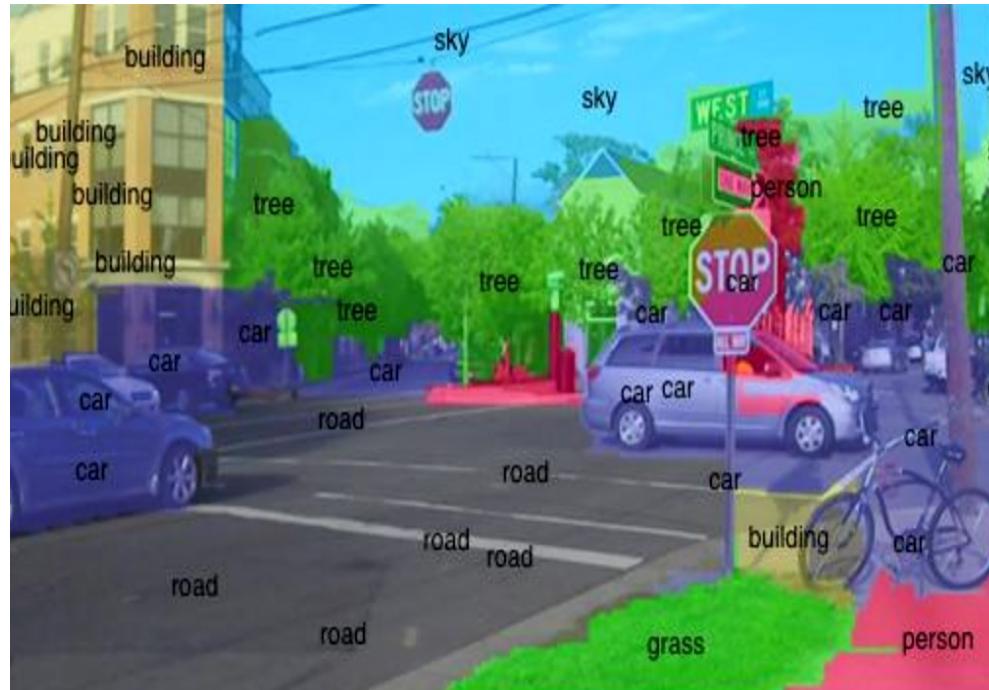


Because Deep Learning delivers state-of-the-art results

Self-Driving Cars



Computer Vision



Natural Language Processing



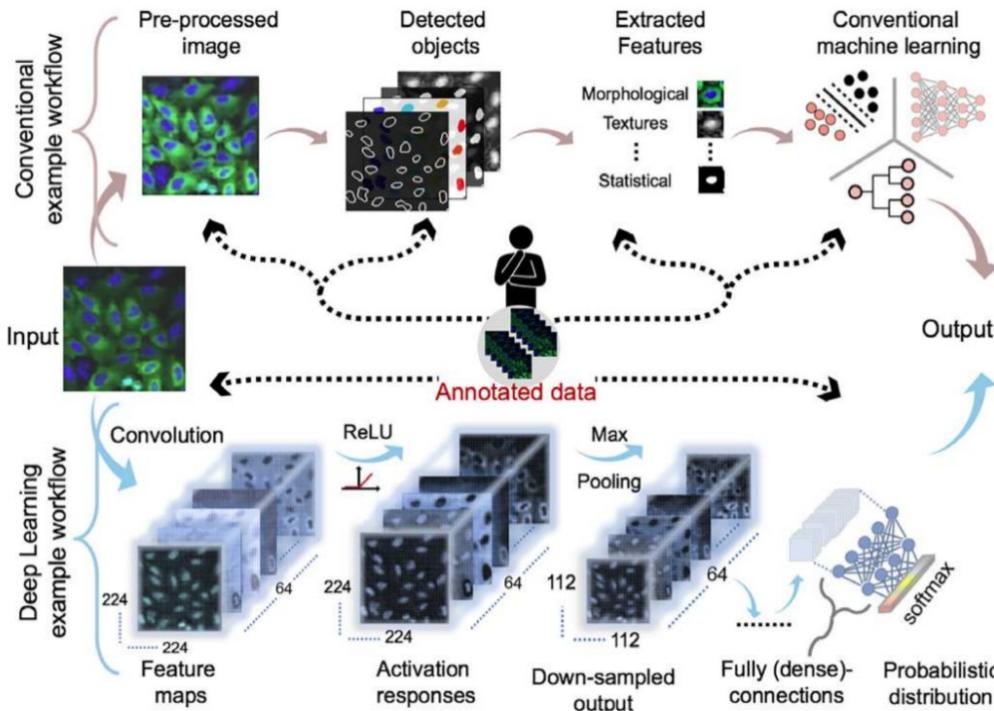
Advanced Search Language Tools	
never put	7,170,000 results
never put a healthy dog down	47,900 results
never put a sock in a toaster	1,220,000 results
never put jam on a magnet	327,000 results
never put a blanket over an owl	130,000,000 results
never put it in writing	62,700 results
never put jam in a toaster	950,000 results
never put new shoes on a table	134,000,000 results
never put on weight	224,000 results
never put your banana in the refrigerator	15,700,000 results
never put baby in a corner	

Speech Recognition





Image Cytometry



Only you can see this message
This story is eligible to be part of the metered payroll. [Learn more](#)

Deep Learning on Ancient DNA

Nikolay Oskolkov
Apr 28 · 6 min read



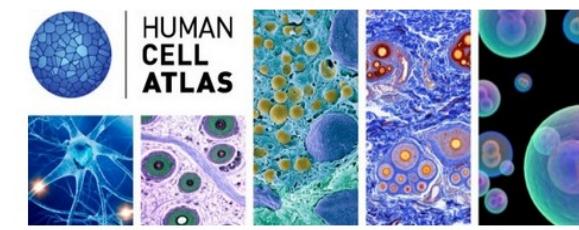
[Image source](#)

Reconstructing the Human Past with Deep Learning

Only you can see this message
This story is eligible to be part of the metered payroll. [Learn more](#)

Deep Learning for Single Cell Biology

Nikolay Oskolkov
May 5 · 9 min read



[Image source](#)

Resolving cellular architectures by Deep Learning

P is the number of features (genes, proteins, genetic variants etc.)

N is the number of observations (samples, cells, nucleotides etc.)

Bayesianism



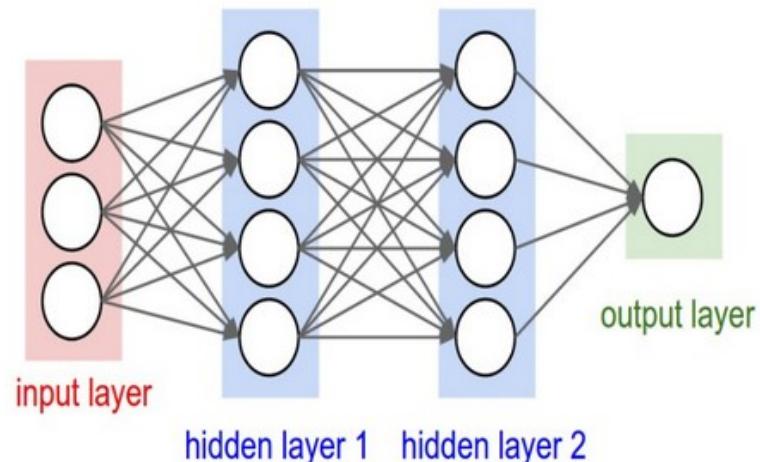
$P \gg N$

Frequentism



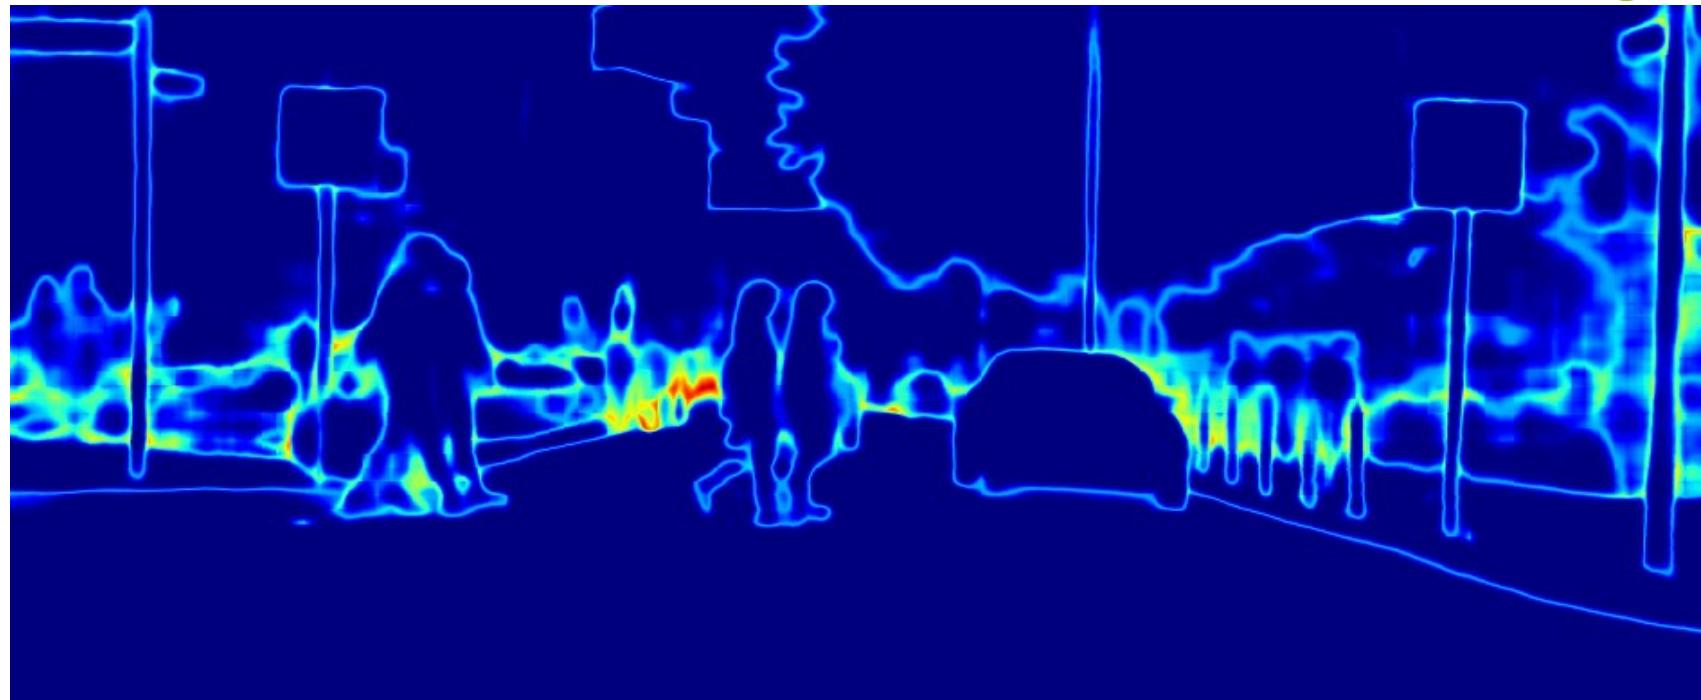
$P \sim N$

Deep Learning

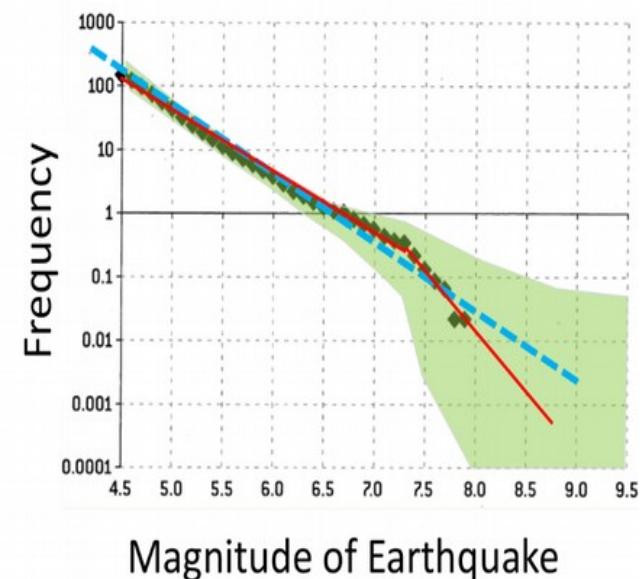
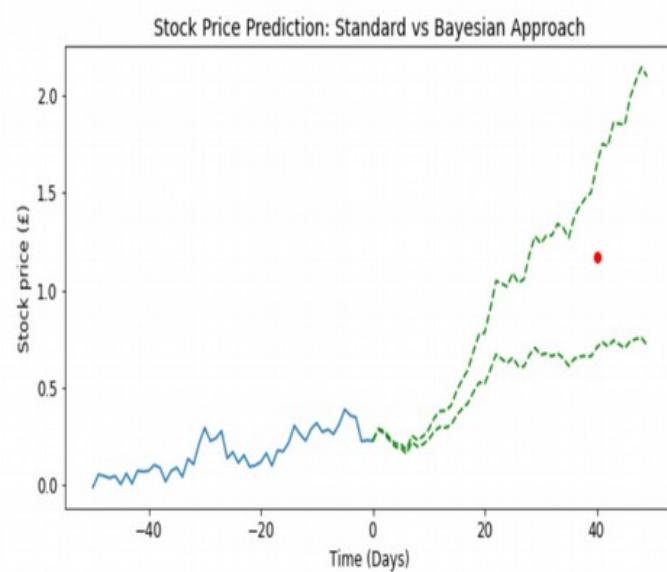


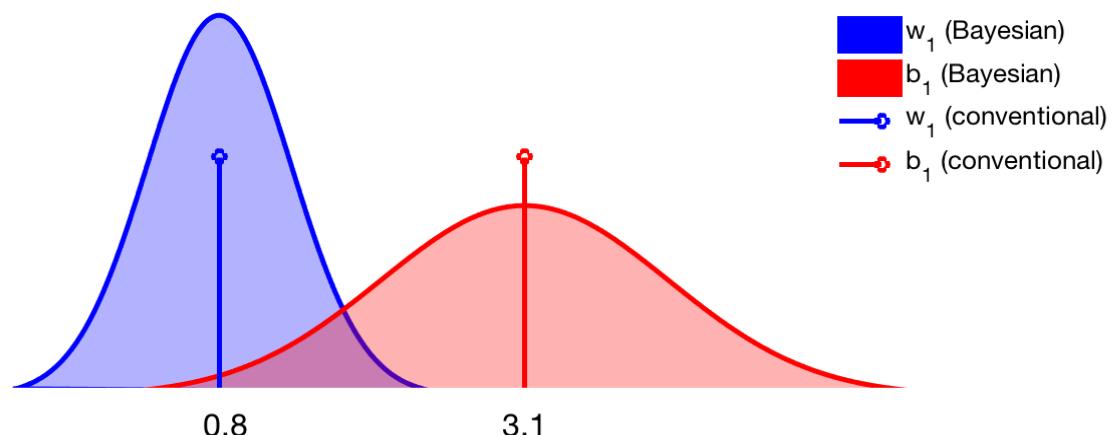
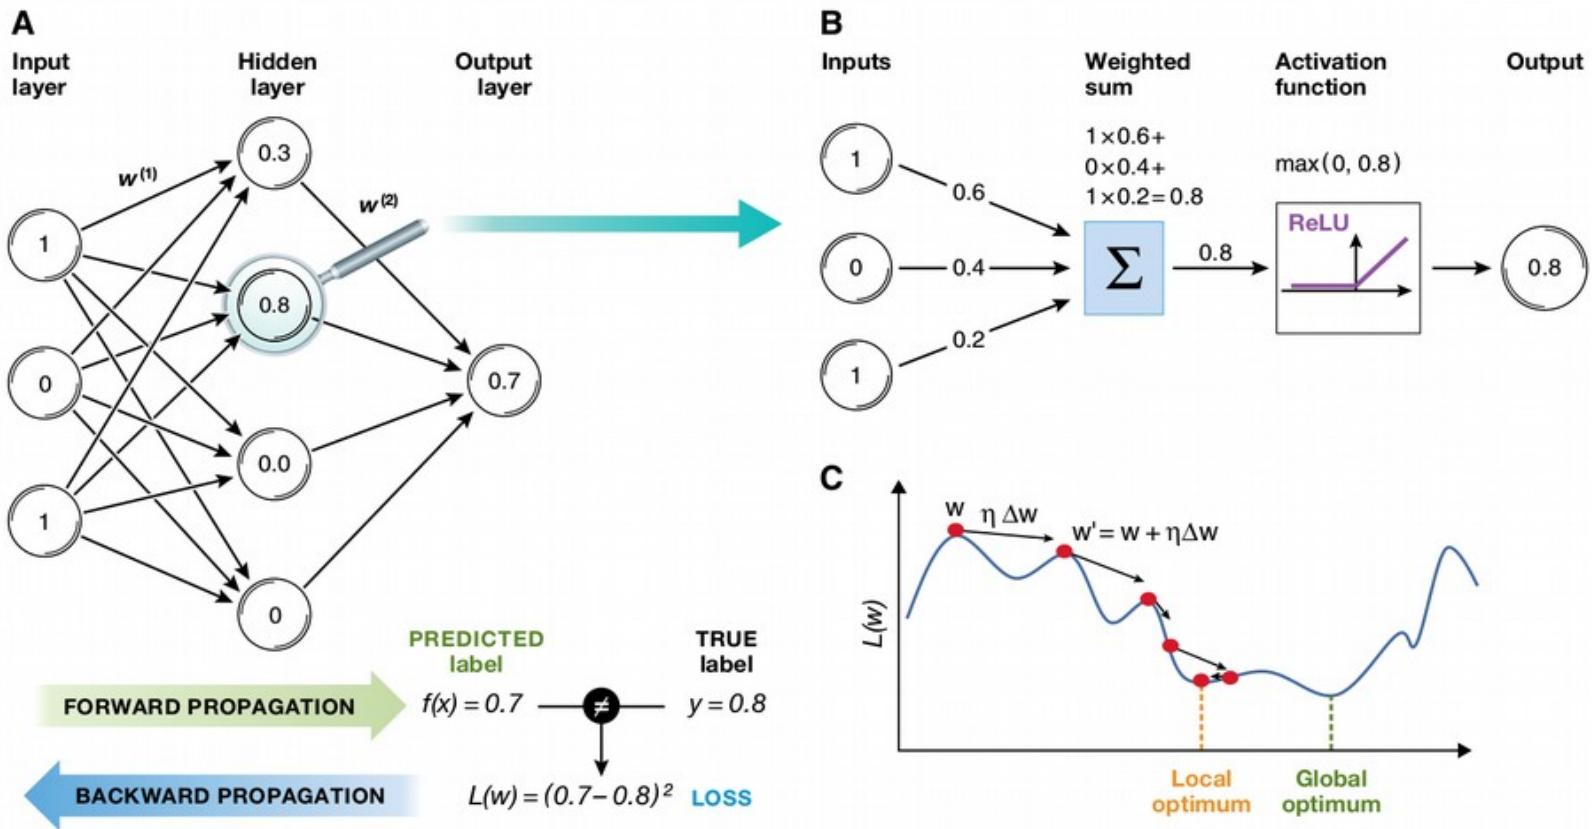
$P \ll N$

Amount of Data



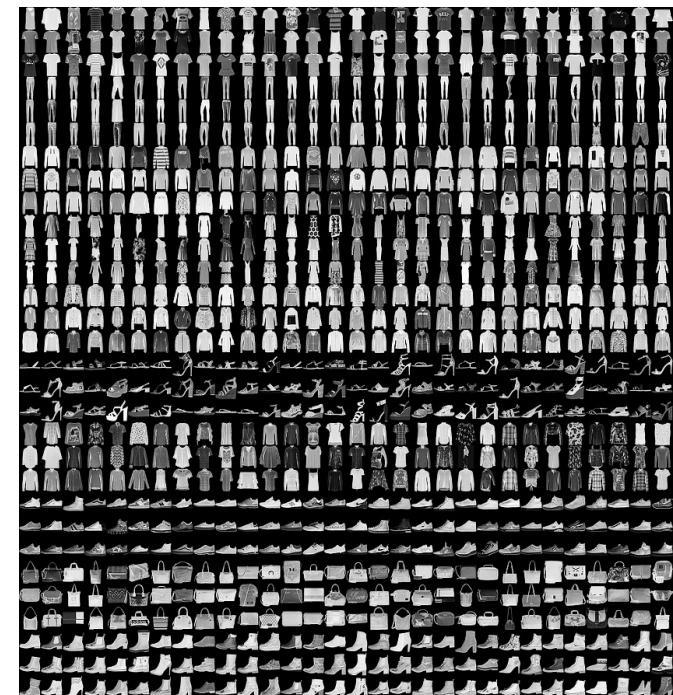
Intelligence is to know how much you do not know





Frequentist Image Recognition

Fashion MNIST



```
In [24]: # normalize inputs from 0-255 to 0-1.0
X_train = X_train.reshape(X_train.shape[0], 1, 28, 28).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 1, 28, 28).astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0

In [25]: # one hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]
print(num_classes)

In [27]: # Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(1, 28, 28), padding='same', activation='relu',
               kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), padding='same', activation='relu',
               kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

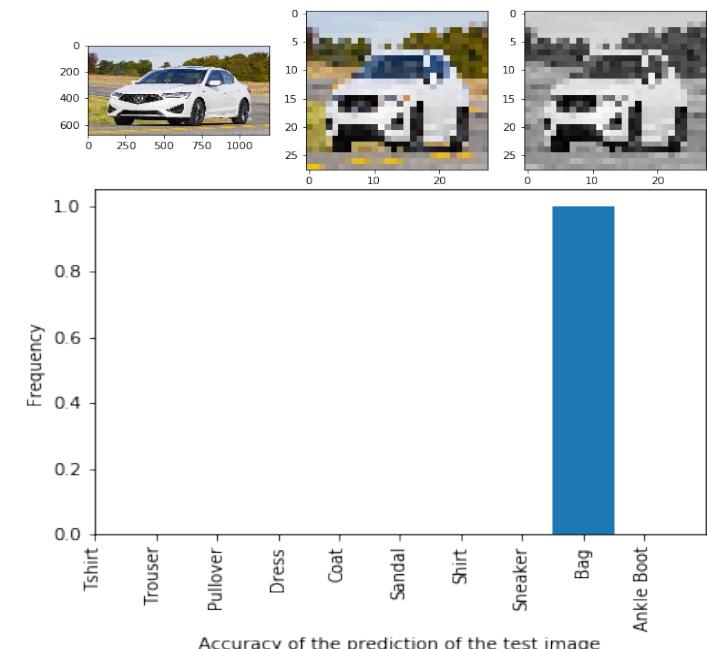
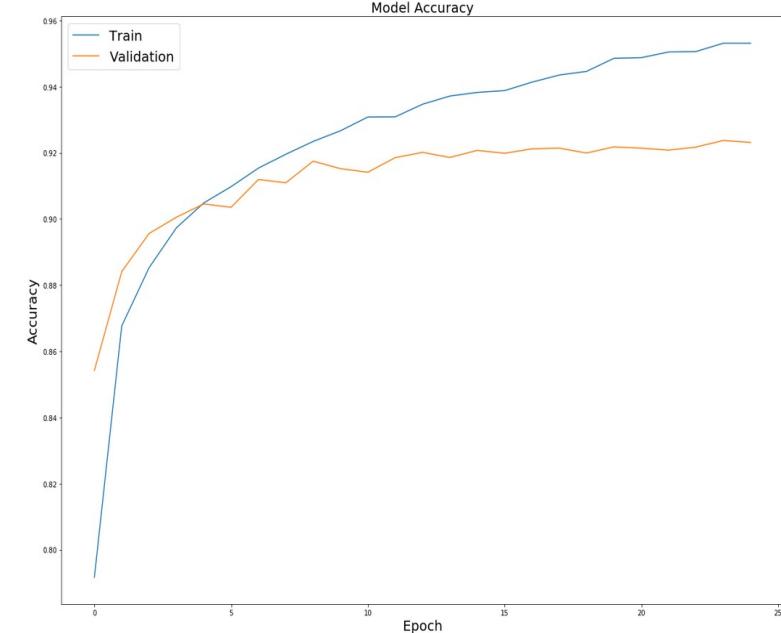
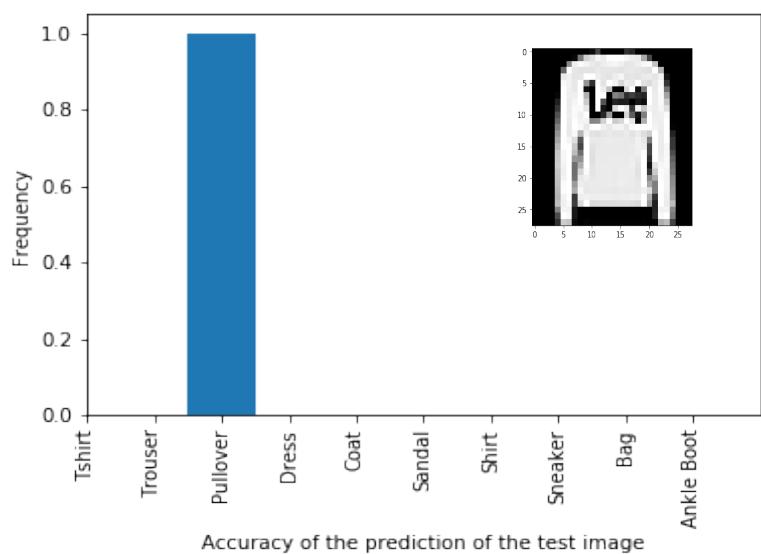
# Compile model
epochs = 25
lr_rate = 0.01
decay = lr_rate/epochs
sgd = SGD(lr=lr_rate, momentum=0.9, decay=decay, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

Layer (type) Output Shape Param #
===== =====
conv2d_8 (Conv2D) (None, 32, 28, 28) 320
dropout_7 (Dropout) (None, 32, 28, 28) 0
conv2d_9 (Conv2D) (None, 32, 28, 28) 9248
max_pooling2d_4 (MaxPooling2D) (None, 32, 14, 14) 0
flatten_4 (Flatten) (None, 6272) 0
dense_7 (Dense) (None, 512) 3211776
dropout_8 (Dropout) (None, 512) 0
dense_8 (Dense) (None, 10) 5130
=====
Total params: 3,226,474
Trainable params: 3,226,474
Non-trainable params: 0
None

In [28]: # Fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32,
          history = model.history, verbose = 1, validation_split = 0.25,
          batch_size = 32, shuffle = True)

Train on 45000 samples, validate on 15000 samples
Epoch 1/25
45000/45000 [=====] - 1158s 26ms/step - loss: 0.5762 - acc: 0.7917 - val_loss: 0.39
73 - val_acc: 0.8542
Epoch 2/25
45000/45000 [=====] - 1124s 25ms/step - loss: 0.3643 - acc: 0.8676 - val_loss: 0.31
27 - val_acc: 0.8841
Epoch 3/25
45000/45000 [=====] - 1158s 26ms/step - loss: 0.3129 - acc: 0.8853 - val_loss: 0.28
25 - val_acc: 0.8956
Epoch 4/25
45000/45000 [=====] - 1069s 36ms/step - loss: 0.2813 - acc: 0.8973 - val_loss: 0.27
27 - val_acc: 0.9095
Epoch 5/25
45000/45000 [=====] - 902s 20ms/step - loss: 0.2618 - acc: 0.9048 - val_loss: 0.258
8 - val_acc: 0.9045
Epoch 6/25
45000/45000 [=====] - 936s 21ms/step - loss: 0.2451 - acc: 0.9098 - val_loss: 0.256
4 - val_acc: 0.9035
Epoch 7/25
```

Prediction



PyMC3, Edward, TensorFlow Probability

```
In [8]: x_train = x_train.reshape((x_train.shape[0],D))
x_test = x_test.reshape((x_test.shape[0],D))
print(x_train.shape)
print(x_test.shape)

(60000, 784)
(10000, 784)

In [9]: from keras.utils import to_categorical
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
print(y_train.shape)
print(y_test.shape)

(60000, 10)
(10000, 10)

In [10]: ed.set_seed(314159)
N = 100 # number of images in a minibatch.
D = D # number of features.
K = 10 # number of classes.

# Create a placeholder to hold the data (in minibatches) in a TensorFlow graph.
x = tf.placeholder(tf.float32, [None, D])
# Normal(0,1) priors for the variables. Note that the syntax assumes TensorFlow 1.1.
w = Normal(loc=tf.zeros([D, K]), scale=tf.ones([D, K]))
b = Normal(loc=tf.zeros(K), scale=tf.ones(K))
# Categorical likelihood for classification.
y = Categorical(tf.matmul(x, w) + b)

In [11]: # Construct the q(w) and q(b). In this case we assume Normal distributions.
qw = Normal(loc=tf.Variable(tf.random_normal([D, K])),
            scale=tf.nn.softplus(tf.Variable(tf.random_normal([D, K]))))
qb = Normal(loc=tf.Variable(tf.random_normal([K])),
            scale=tf.nn.softplus(tf.Variable(tf.random_normal([K]))))

In [12]: def generator(arrays, batch_size = N):
    starts = [0] * len(arrays) # pointers to where we are in iteration
    while True:
        batches = []
        for i, array in enumerate(arrays):
            start = starts[i]
            stop = start + batch_size
            diff = stop - array.shape[0]
            if diff <= 0:
                batch = array[start:stop]
                starts[i] += batch_size
            else:
                batch = np.concatenate((array[start:], array[:diff]))
                starts[i] = diff
            batches.append(batch)
        yield batches
    cifar10 = generator([x_train, y_train], N)

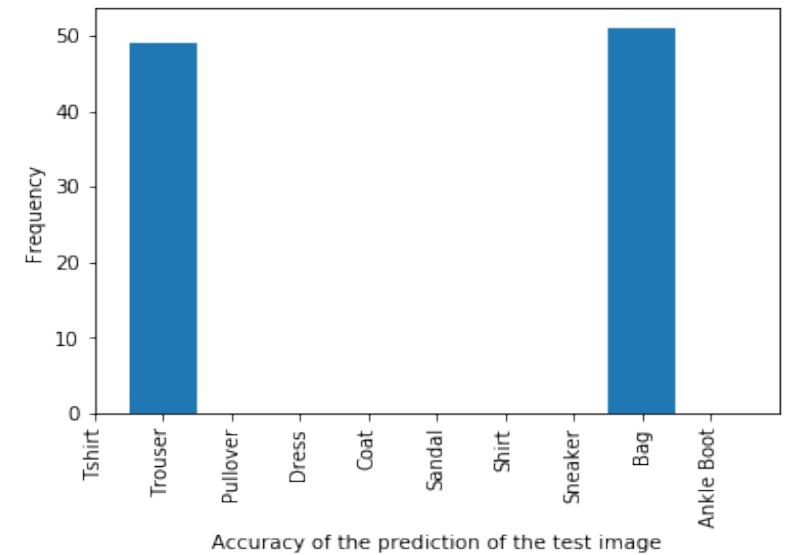
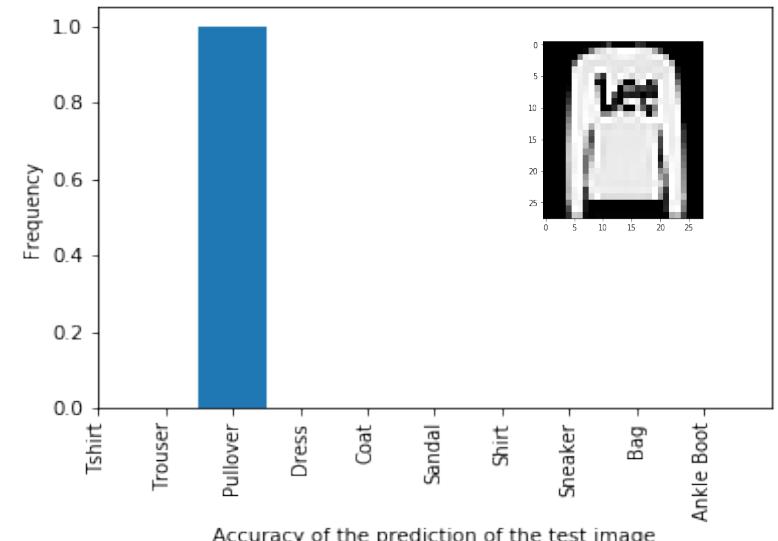
In [13]: # We use a placeholder for the labels in anticipation of the training data.
y_ph = tf.placeholder(tf.int32, [None])
# Define the VI inference technique, ie. minimise the KL divergence between q and p.
inference = ed.KLqp(qw, {w: qb}, data={y: y_ph})
# Initialise the session.
sess = tf.InteractiveSession()
# Initialise all the variables in the session.
tf.global_variables_initializer().run()
# Let the training begin. We load the data in minibatches and update the VI inference using each new batch.
for _ in range(inference.n_iter):
    X_batch, Y_batch = next(cifar10)
    #X_batch = X_batch.reshape(-1, 1)
    #TensorFlow method gives the label data in a one hot vector format. We convert that into a single label.
    Y_batch = np.argmax(Y_batch, axis=1)
    info_dict = inference.update(feed_dict={x: X_batch, y_ph: Y_batch})
    inference.print_progress(info_dict)
50000/50000 [100%] Elapsed: 221s | Loss: 85453.266

In [14]: # Generate samples the posterior and store them.
n_samples = 100
prob_lst = []
samples = []
w_samples = []
b_samples = []
for _ in range(n_samples):
    w_samp = qw.sample()
    b_samp = qb.sample()
    w_samples.append(w_samp)
    b_samples.append(b_samp)
    # Also compute the probability of each class for each (w,b) sample.
    prob = tf.nn.softmax(tf.matmul(x_test, w_samp) + b_samp)
    prob_lst.append(prob.eval())
    sample = tf.concat([tf.reshape(w_samp, [-1]), b_samp], 0)
    samples.append(sample.eval())

In [15]: # Compute the accuracy of the model.
# For each sample we compute the predicted class and compare with the test labels.
# Predicted class is defined as the one which has maximum probability.
# We perform this test for each (w,b) in the posterior giving us a set of accuracies
# Finally we take a histogram of accuracies for the test data.
dict_acccy = {}
for prob in prob_lst:
    y_trn_prd = np.argmax(prob, axis=1).astype(np.float32)
    acc = (y_trn_prd == np.argmax(y_test, axis=1)).mean()*100
    accy_test.append(acc)

plt.hist(acccy_test)
plt.title("Histogram of prediction accuracies in the CIFAR10 test data")
plt.xlabel("Accuracy")
plt.ylabel("Frequency")
plt.show()
```

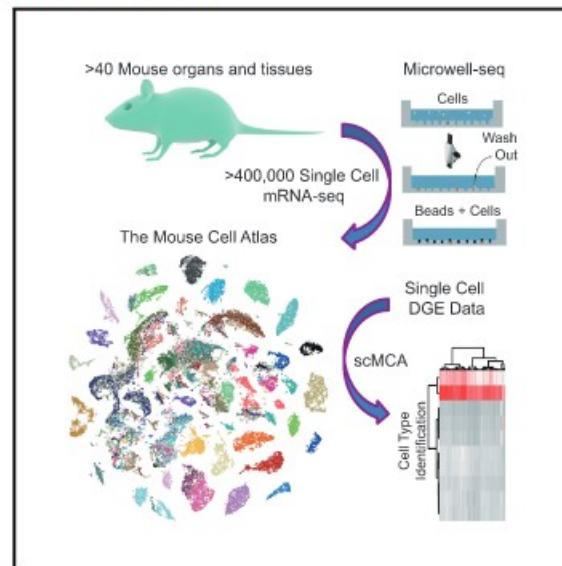
Prediction



Cell

Mapping the Mouse Cell Atlas by Microwell-Seq

Graphical Abstract



Authors

Xiaoping Han, Renying Wang,
Yincong Zhou, ..., Guo-Cheng Yuan,
Ming Chen, Guoji Guo

Correspondence

xhan@zju.edu.cn (X.H.),
ggj@zju.edu.cn (G.G.)

In Brief

Development of Microwell-seq allows construction of a mouse cell atlas at the single-cell level with a high-throughput and low-cost platform.

Data Resources

GSE108097

Highlights

- Development of Microwell-seq, a high-throughput and low-cost scRNA-seq platform
- Construction of a single-cell mouse cell atlas (scMCA) covering major cell types
- Characterization of cellular heterogeneity with minimal batch effect
- Characterization of cross-tissue cellular network at the single-cell level

Resource

CAREERS BLOG 10X UNIVERSITY EV

10X GENOMICS

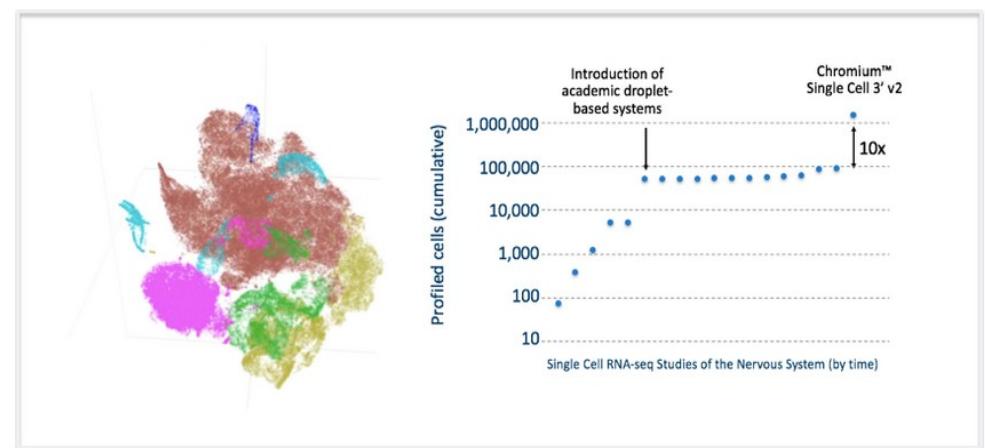
SOLUTIONS & PRODUCTS

RESEARCH & APPLICATIONS

EDUCATION & RESOURCES

[« Back to Blog](#)

[« Newer Article](#) [Older Article »](#)



Our 1.3 million single cell dataset is ready 0 KUDOS

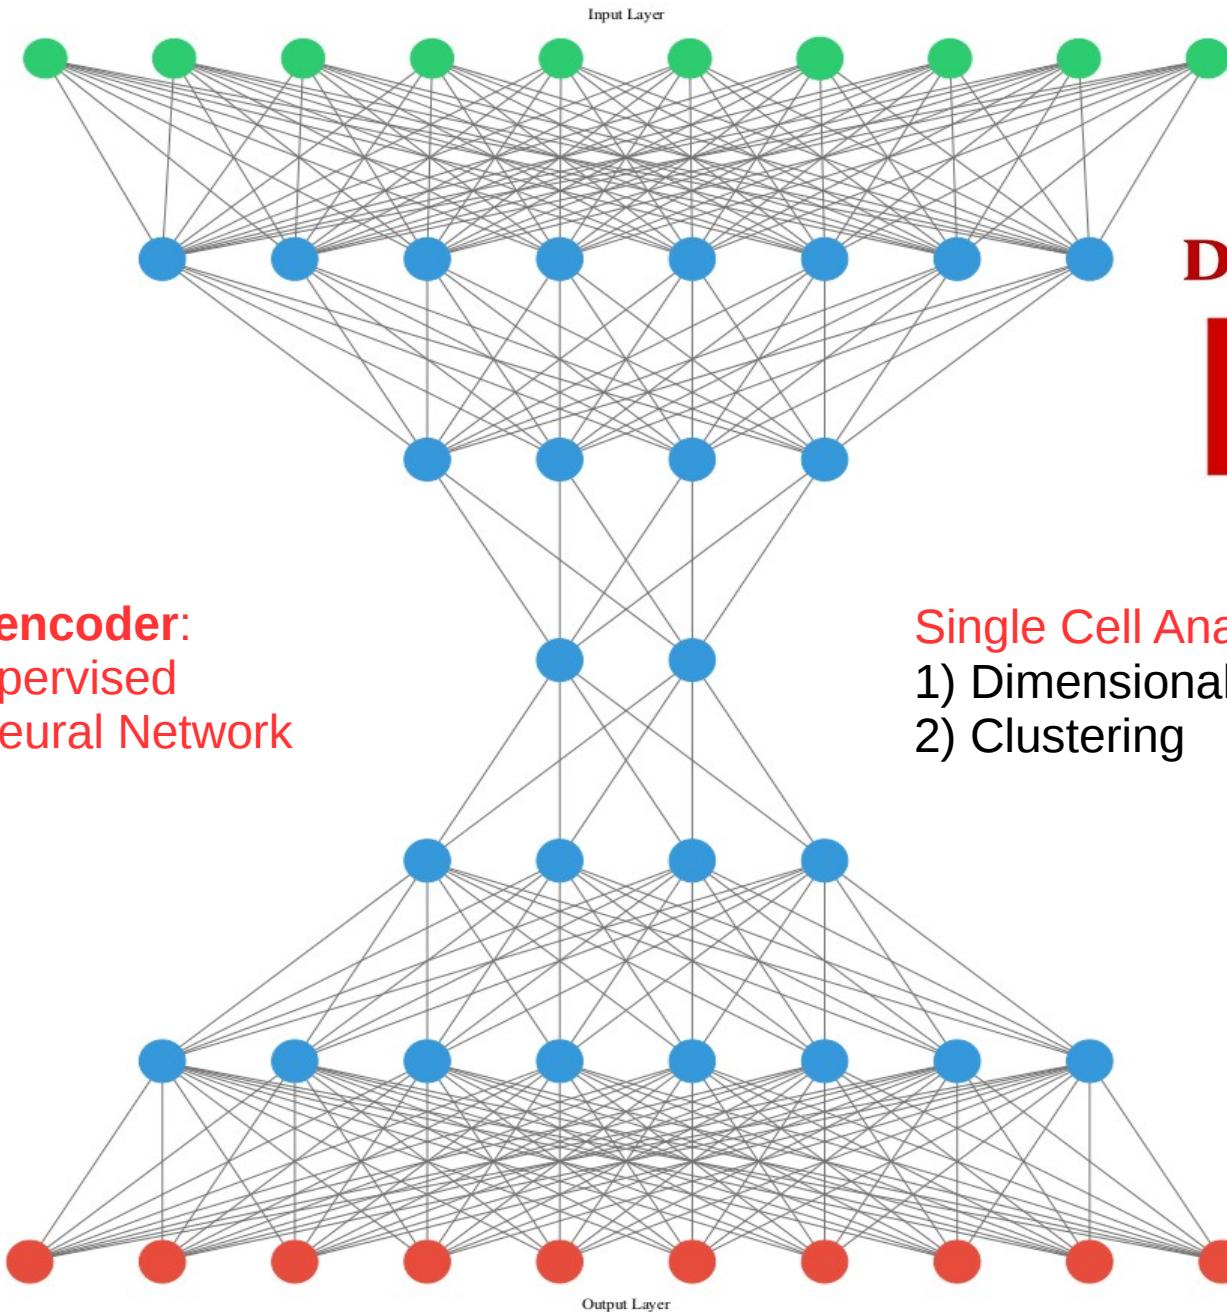


POSTED BY grace-10x, on Feb 21, 2017 at 2:28 PM

At ASHG last year, we announced our 1.3 Million Brain Cell Dataset, which is, to date, the largest dataset published in the single cell RNA-sequencing (scRNA-seq) field. Using the Chromium™ Single Cell 3' Solution (v2 Chemistry), we were able to sequence and profile 1,308,421 individual cells from embryonic mice brains. Read more in our application note [Transcriptional Profiling of 1.3 Million Brain Cells with the Chromium™ Single Cell 3' Solution](#).



Han et al., 2018, Cell 172, 1091–1107
February 22, 2018 © 2018 Elsevier Inc.
<https://doi.org/10.1016/j.cell.2018.02.001>

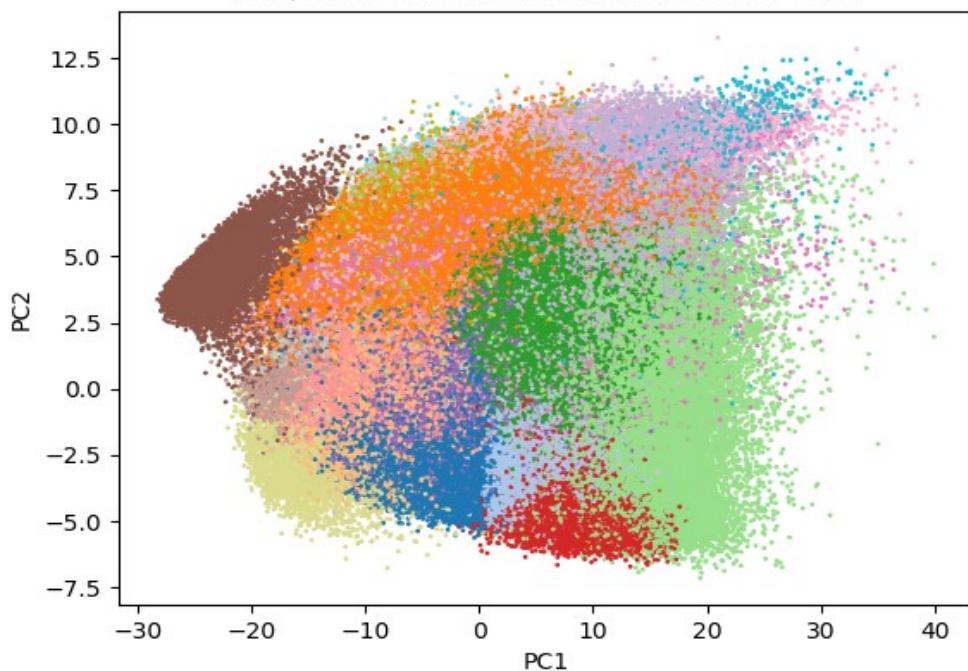


Autoencoder:
Unsupervised
Artificial Neural Network

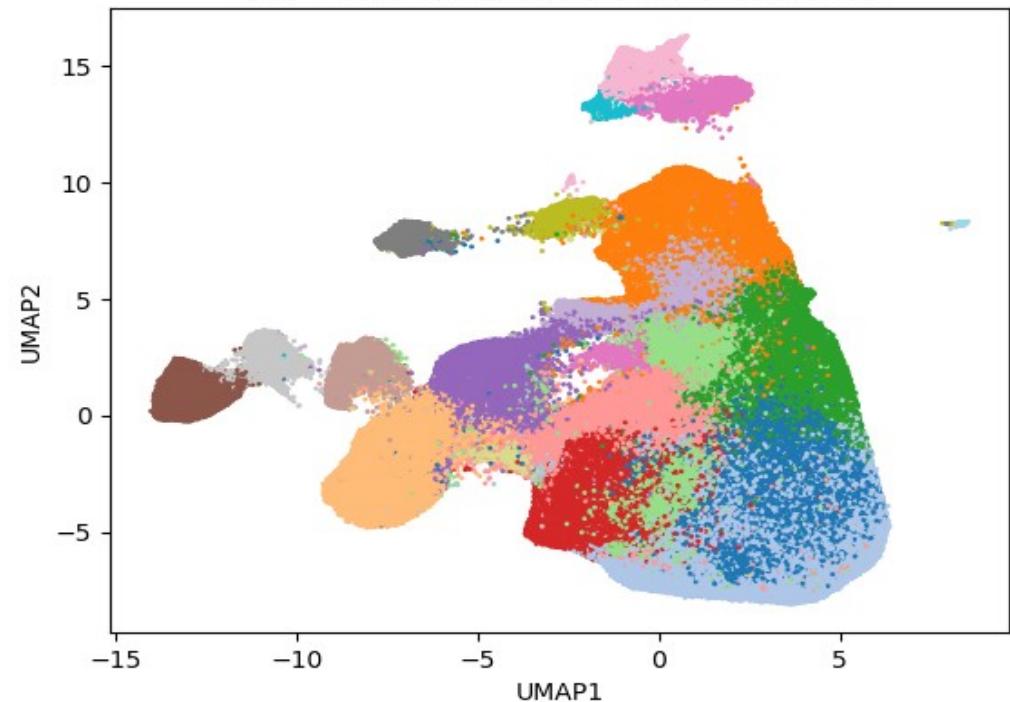
Single Cell Analysis is Unsupervised
1) Dimensionality Reduction
2) Clustering

Deep Learning
K + TensorFlow

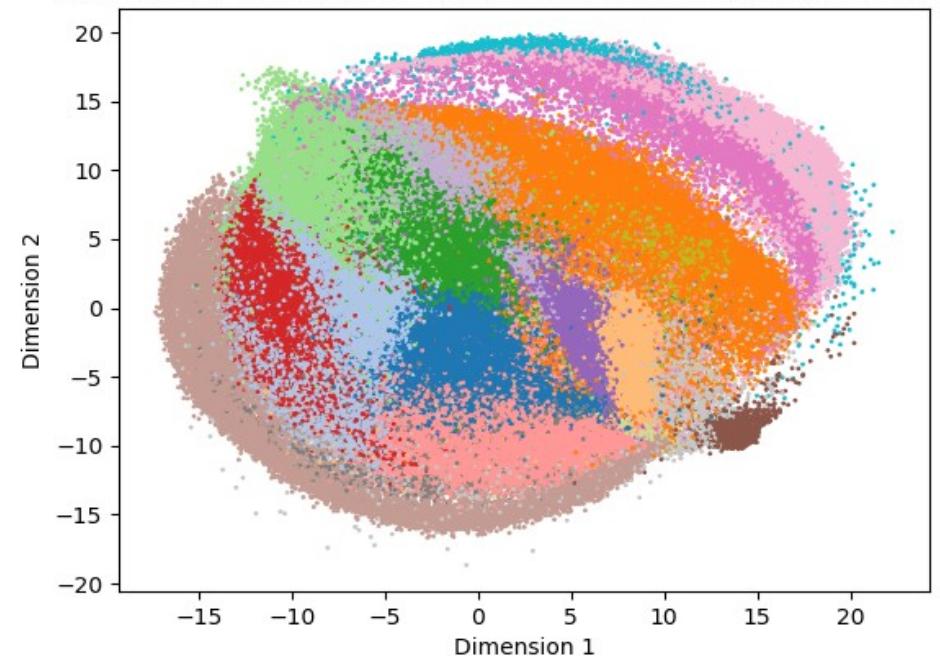
PCA, 10X Genomics 1.3M Mouse Brain Cells



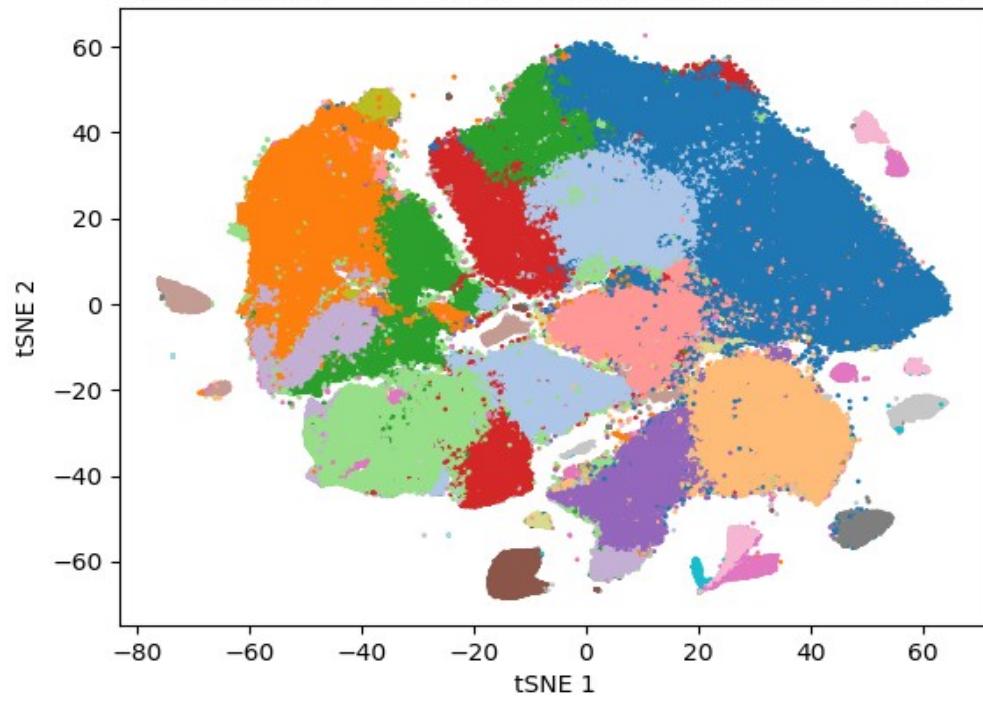
UMAP 10X Genomics 1.3M Mouse Brain cells

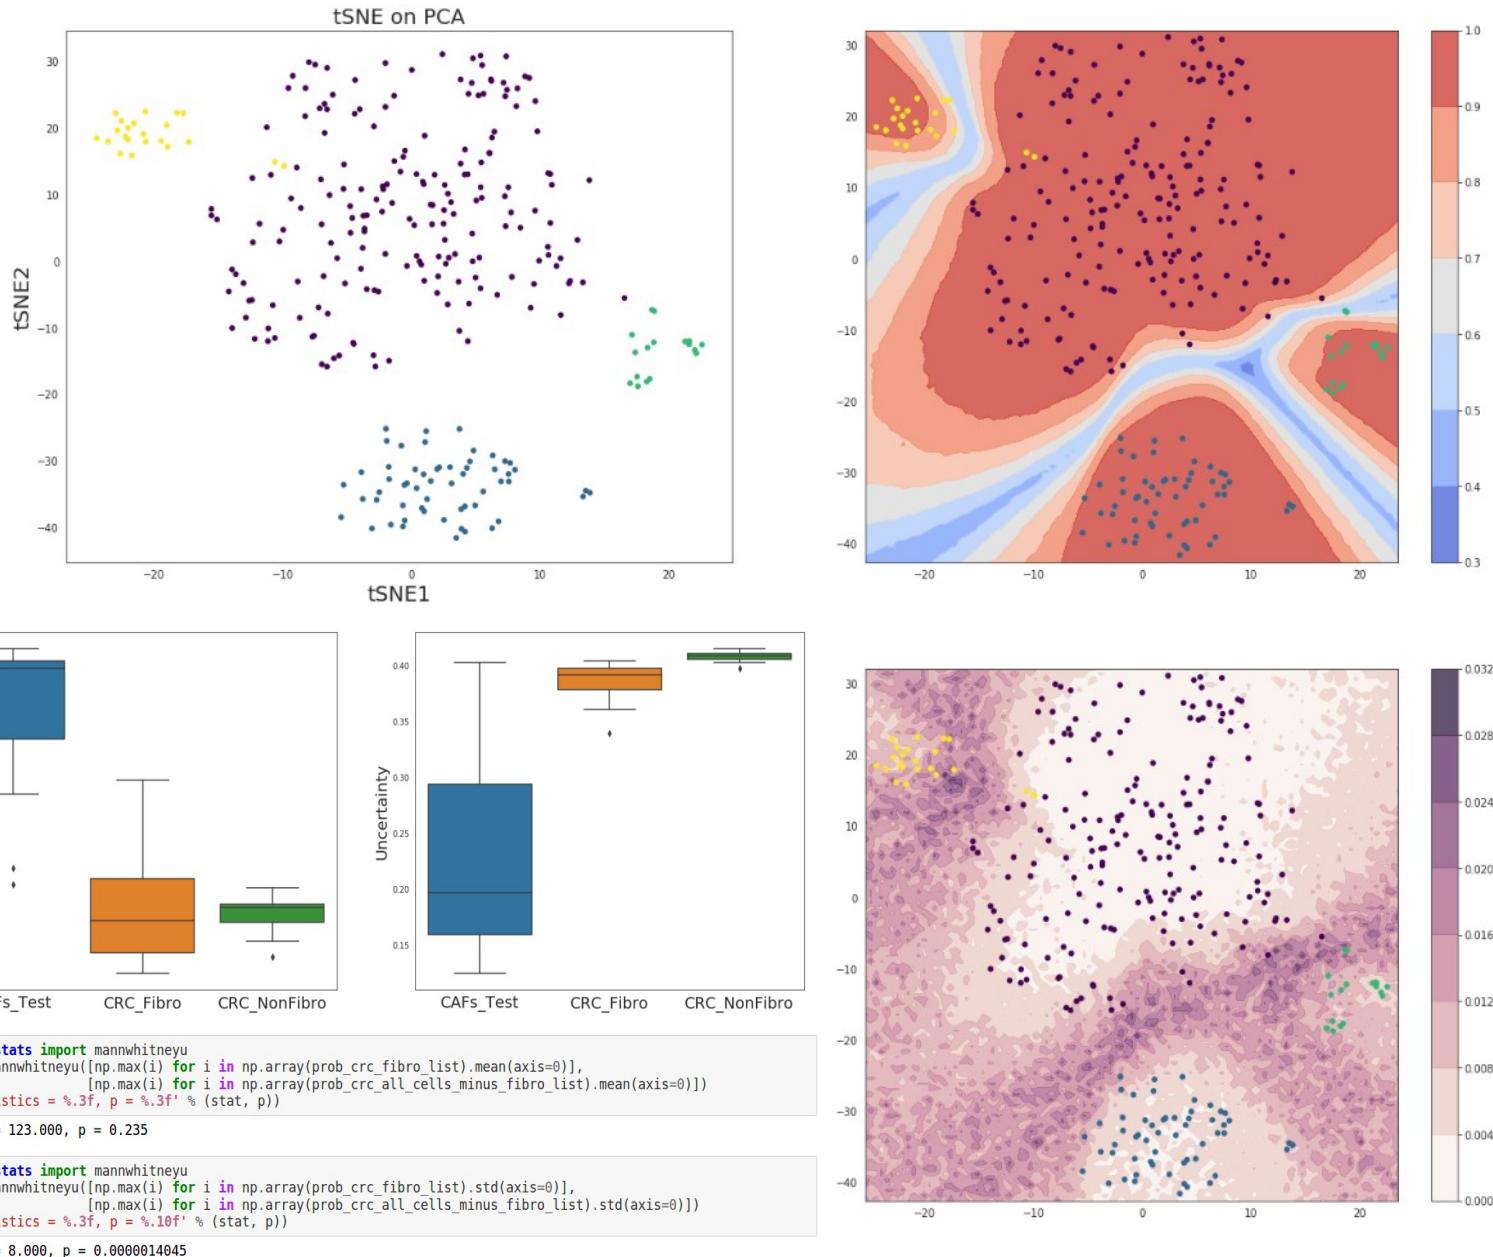


Autoencoder 10 Hiden Layers, 10X Genomics 1.3M Mouse Brain cells



tSNE perplexity = 350, 10X Genomics 1.3M Mouse Brain cells





Bartoschek et al. 2018, Nature Communications, 9, 5150

Superior for predictions on unseen data



National Bioinformatics Infrastructure Sweden (NBIS)

SciLifeLab



*Knut och Alice
Wallenbergs
Stiftelse*



LUNDS
UNIVERSITET